

Differential Dynamic Programming

Ramona Stefanescu

October 18, 2018

MDPs

- ▶ A (discounted infinite horizon) Markov Decision Process (MDP) is a tuple (S, A, T, γ, D, R)
- ▶ S is the set of possible states for the system
- ▶ A is the set of possible actions
- ▶ T represents the (typically stochastic) system dynamics
- ▶ D is the initial-state distribution, from which state s_0 is drawn
- ▶ $R : S \rightarrow R$ is the reward function

Acting in a Markov decision process results in a sequence of states and actions $s_0, a_0, s_1, s_2, \dots$

A policy π is a sequence of mappings $(\mu_0, \mu_1, \mu_2 \dots)$, where, at time t the mapping $\mu_t(\cdot)$ determines the action $a_t = \mu_t(s_t)$ to take when in state s_t

- ▶ The objective is to find policies that maximize the expected sum of rewards accumulated over time
- ▶ In particular, a policy π is good if its utility is high:

$$U(\pi) = \mathbb{E}\left[\sum_{t=0}^{\infty} \gamma^t R(s_t) | \pi\right] \quad (1)$$

- ▶ To represent the system dynamics, we can use the state-transition distribution notation

$$s_{t+1} \propto P_{sa}(\cdot | s_t, a_t) \quad (2)$$

- ▶ Or using a deterministic function F and a random disturbance ω_t

$$s_{t+1} = F(s_t, a_t, \omega_t) \quad (3)$$

Example - car

- ▶ One way to model the state of a car is to use the following six state variables: northing (n), easting (e), north velocity (\dot{n}), east velocity (\dot{e}), heading (θ), angular rate ($\dot{\theta}$). Hence the state space $\mathbb{S} = \mathbb{R}^6$
- ▶ The actions (or control inputs) are (i) steering angle, (ii) throttle, (iii) brake
- ▶ The perturbances capture both environmental perturbations as well as unmodeled aspects of the car dynamics
- ▶ We could have the following dynamics model $s_{t+1} = F(s_t, a_t, \omega_t)$:

$$n_{t+1} = n_t + \dot{n}_t \Delta_t, \quad (4)$$

$$e_{t+1} = e_t + \dot{e}_t \Delta_t, \quad (5)$$

$$\theta_{t+1} = \theta_t + \dot{\theta}_t \Delta_t \quad (6)$$

$$\dot{n}_{t+1} = f_n(\dot{n}_t, \dot{e}_t, \dot{\theta}_t, a_t, \omega_t) \quad (7)$$

$$\dot{e}_{t+1} = f_e(\dot{n}_t, \dot{e}_t, \dot{\theta}_t, a_t, \omega_t) \quad (8)$$

$$\dot{\theta}_{t+1} = f_\theta(\dot{n}_t, \dot{e}_t, \dot{\theta}_t, a_t, \omega_t) \quad (9)$$

- ▶ The reward function could be $R(s_t) = 1$ (in goal region) and 100 (in collision)

Finding optimal policies: value iteration

The goal is to find an optimal policy - the policy that maximizes the expected total of rewards earned over the period of our decision process.

- ▶ **Finite horizon:** we are concerned with decisions and rewards only up until a given time $t = H$, with state space S and action space A finite.
- ▶ The value of any policy π is:

$$V_{\pi}(s_0) = \mathbb{E}\left[\sum_{t=0}^H \gamma^t R(s_t) | \pi; s_0\right] \quad (10)$$

we are interested in finding

$$\max_{\pi} V_{\pi}(s_0) \quad (11)$$

$$\pi = \{\mu_0, \mu_1, \dots, \mu_{H-1}\} \quad (12)$$

where $\mu_i : S \rightarrow A$

Finite horizon

- ▶ Since there are $|A|^{|S|}$ possible mappings for each μ_i , there are $(|A|^{|S|})^H$ possible policies π , which is far too many to compute the value of all possible options
- ▶ Instead, we apply a dynamic programming algorithm known as value iteration to find the optimal policy efficiently
- ▶ Intuitively, we are applying the notion that given a state, the past and future are independent (the "Markov property")
- ▶ Define the value function at the k th time-step as

$$V_k(s_k) = \max_{\mu_k, \mu_{k+1}, \dots, \mu_{H-1}} \mathbb{E} \left[\sum_{t=k}^H \gamma^{t-k} R(s_t) | s_k \right] \quad (13)$$

- ▶ It can be shown that the optimal policy can be found by working backwards from H :

$$V_k(s_k) = \max_{a \in A} [R(s_k) + \gamma \sum_{s'} P(s' | s, a) V_{k+1}(s')] \quad (14)$$

Infinite horizon optimal policies

- ▶ Now we are interested in finding an optimal policy when our horizon is infinite
- ▶ The value of a policy is now

$$V_{\pi}(s_0) = \mathbb{E}\left[\sum_{t=0}^{\infty} \gamma^t R(s_t) | s_0, \pi\right] \quad (15)$$

- ▶ And we are interested in finding $V^*(s) = \max_{\pi} V_{\pi}(s)$ and $\pi^* = \operatorname{argmax}_{\pi} V_{\pi}(s)$
- ▶ Despite the fact the expectation is taken over an infinite sum, it is guaranteed to converge. This is because R is a function over a finite state space, and is therefore bounded, and since $\gamma \in (0, 1)$

Bellman Backup Operator

- ▶ Let $V : S \rightarrow R$ be our value function and let $V_\pi(s) = \mathbb{E}[\sum_{t=0}^{\infty} \gamma^t R(s_t) | s_0, \pi]$.
- ▶ Define the operator $T : V \leftarrow TV$ as

$$(TV)(s) = \max_{a \in A} [R(s) + \gamma \sum_{s'} P(s'|s, a) V(s')] \quad (16)$$

LQR - Finite Horizon Value Iteration Case Study

- ▶ We assume a linear dynamic system:

$$x_{t+1} = A_t x_t + B_t u_t, \quad (17)$$

where $x(t)$ denotes the state at time t and $u(t)$ denotes the input at time t

- ▶ We assume a quadratic cost function:

$$J = \sum_{t=0}^{H-1} (x_t^T Q_t x_t + u_t^T R_t u_t) + x_H^T P_H x_H \quad (18)$$

with for all t , $Q_t \geq 0, R_t \geq 0$

- ▶ Our goal is to find the input sequence $\{u_0, u_1, \dots, u_{H-1}\}$ that minimizes the cost: $\min_{u_0 \dots u_{H-1}} J$

LQR - cont

$$\begin{aligned} & \min_{u_0 \dots u_{H-1}} J \\ &= \min_{u_0 \dots u_{H-1}} \sum_{t=0}^{H-1} (x_t^T Q_t x_t + u_t^T R_t u_t) + x_H^T P_H x_H \\ &= \min_{u_0 \dots u_{H-2}} \sum_{t=0}^{H-2} (x_t^T Q_t x_t + u_t^T R_t u_t) + x_{H-1}^T Q_{H-1} x_{H-1} + \\ &+ \min_{u_{H-1}} u_{H-1}^T R_{H-1} u_{H-1} + x_H^T Q_H x_H \end{aligned}$$

LQR - cont

Solving the equation the previous page leads to the following dynamic programming algorithm to find the optimal controller for a linear system with quadratic costs:

- ▶ for $t = H - 1$ to 0
- ▶ $\rightarrow K_t \leftarrow -(R_t + B_t^T P_{t+1} B_t)^{-1} B_t^T P_{t+1} A_t$
- ▶ $\rightarrow P_t \leftarrow Q_t + K_t^T R_t K_t + (A_t + B_t K_t)^T P_{t+1} (A_t + B_t K_t)$
- ▶ next t

The optimal inputs are computed as follows $u_t = K_t x_t$. The optimal cost-to-go (=cost incurred in all future steps) for being in state x_t at time t is given by $x_t^T P_t x_t$.

Controllability vs Observability

